
Problem Solving Agents & Problem Formulation

AIMA 2.3, 3.1-3

Outline for today's lecture

- ***Defining Task Environments (AIMA 2.3)***
- **Environment types**
- **Formulating Search Problems**
- **Search Fundamentals**



Task environments

- To design a rational agent we need to specify a *task environment*
 - a problem specification for which the agent is a solution
- **PEAS:** to specify a task environment
 - **P**erformance measure
 - **E**nvironment
 - **A**ctuators
 - **S**ensors



PEAS: Specifying an automated taxi driver

Performance measure:

- ?

Environment:


- ?

Actuators:

- ?

Sensors:

- ?



PEAS: Specifying an automated taxi driver

Performance measure:

- safe, fast, legal, comfortable, maximize profits

Environment:


- roads, other traffic, pedestrians, customers

Actuators:

- steering, accelerator, brake, signal, horn

Sensors:

- cameras, sonar, speedometer, GPS



PEAS: Medical diagnosis system

- ***P*erformance measure:** Healthy patient, minimize costs, lawsuits
- ***E*nvironment:** Patient, hospital, staff
- ***A*ctuators:** Screen display (form including: questions, tests, diagnoses, treatments, referrals)
- ***S*ensors:** Keyboard (entry of symptoms, findings, patient's answers)

Outline for today's lecture

- Defining Task Environments
- *Environment types (also ALMA 2.3)*
- Formulating Search Problems
- Search Fundamentals

Environment types: Definitions I

- ***Fully observable*** (vs. partially observable): An agent's sensors give it access to the complete state of the environment at each point in time.
- ***Deterministic*** (vs. stochastic): The next state of the environment is completely determined by the current state and the action executed by the agent.
 - If the environment is deterministic except for the actions of other agents, then the environment is ***strategic***.
- ***Episodic*** (vs. sequential): The agent's experience is divided into atomic "episodes" during which the agent perceives and then performs a single action, and the choice of action in each episode depends only on the episode itself.

Environment types: Definitions II

- **Static** (vs. dynamic): The environment is unchanged while an agent is deliberating.
 - The environment is **semidynamic** if the environment itself does not change with the passage of time but the agent's performance score does.
- **Discrete** (vs. continuous): A limited number of distinct, clearly defined percepts and actions.
- **Single agent** (vs. multiagent): An agent operating by itself in an environment.

(See examples in AIIMA, however I don't agree with some of the judgments)

Environment Restrictions for Now

- We will assume environment is
 - *Static*
 - *Fully Observable*
 - *Deterministic*
 - *Discrete*

The rational agent designer's goal

- Goal of AI practitioner who designs rational agents:
given a *PEAS* task environment,
 1. Construct agent function f that maximizes (the expected value of) the performance measure,
 2. Design an agent program that implements f on a particular architecture

Outline for today's lecture

- Defining Task Environments
- Environment types
- *Formulating Search Problems (AIMA, 3.1-3.2)*
- Search Fundamentals

Example search problem: 8-puzzle



- Formulate *goal*

- Pieces to end up in order as shown...

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- Formulate *search problem*

- **States:** configurations of the puzzle (9! configurations)
- **Actions:** Move one of the movable pieces (≤ 4 possible)
- **Performance measure:** minimize total moves

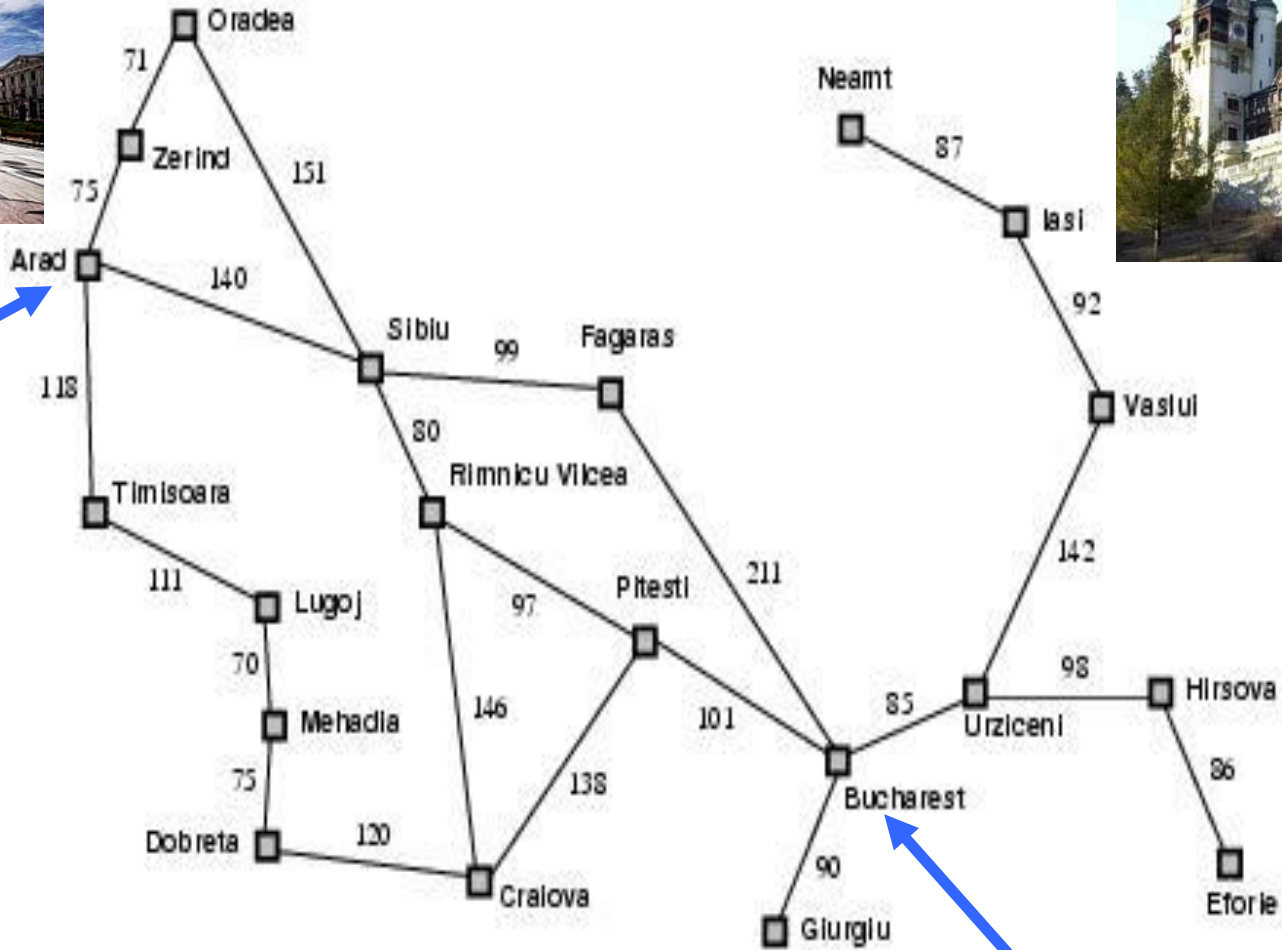
- Find *solution*

- Sequence of pieces moved: 3,1,6,3,1,...

Example search problem: holiday in Romania



You are here



You need to
be here

Holiday in Romania II

- **On holiday in Romania; currently in Arad**
 - Flight leaves tomorrow from Bucharest
- **Formulate *goal***
 - Be in Bucharest
- **Formulate *search problem***
 - States: various cities
 - Actions: drive between cities
 - Performance measure: minimize distance
- **Find *solution***
 - Sequence of cities; e.g. Arad, Sibiu, Fagaras, Bucharest,
...

More formally, a problem is defined by:

1. A set of *states* S
2. An *initial state* $s_i \in S$
3. A *set of actions* A
 - $\forall s \text{ Actions}(s) = \text{the set of actions that can be executed in } s, \text{ that are applicable in } s.$
4. *Transition Model*: $\forall s \forall a \in \text{Actions}(s) \text{ Result}(s, a) \rightarrow s_r$
 - s_r is called a *successor* of s
 - $\{s_i\} \cup \text{Successors}(s_i)^* = \text{state space}$
5. *Goal test* $\text{Goal}(s)$
 - Can be implicit, e.g. *checkmate*(x)
 - s is a *goal state* if $\text{Goal}(s)$ is *true*
6. *Path cost* (additive)
 - e.g. sum of distances, number of actions executed, ...
 - $c(x, a, y)$ is the step cost, assumed ≥ 0
 - (where action a goes from state x to state y)

Solution

A ***solution*** is a sequence of actions from the ***initial state*** to a ***goal state***.

Optimal Solution:

A solution is ***optimal*** if no solution has a lower path cost.

Hard subtask: Selecting a state space

- **Real world is absurdly complex**
State space must be *abstracted* for problem solving
- **(abstract) *State* = set (equivalence class) of real world states**
- **(abstract) *Action* = equivalence class of combinations of real world actions**
 - e.g. *Arad* → *Zerind* represents a complex set of possible routes, detours, rest stops, etc
 - The abstraction is valid if the path between two states is reflected in the real world
- **Each abstract action should be “easier” than the real problem**

Art: Formulating a Search Problem

Decide:

- **Which properties matter & how to represent**
 - *Initial State, Goal State, Possible Intermediate States*
- **Which actions are possible & how to represent**
 - *Operator Set: Actions and Transition Model*
- **Which action is next**
 - *Path Cost Function*

Formulation greatly affects combinatorics of search space and therefore speed of search

Example: 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- States??
- Initial state??
- Actions??
- Transition Model??
- Goal test??
- Path cost??

Example: 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- **States??** **List of 9 locations- e.g., [7,2,4,5,-,6,8,3,1]**
- **Initial state??** **[7,2,4,5,-,6,8,3,1]**
- **Actions??** **{Left, Right, Up, Down}**
- **Transition Model?? ...**
- **Goal test??** **Check if goal configuration is reached**
- **Path cost??** **Number of actions to reach goal**

Example: Missionaries & Cannibals

Three missionaries and three cannibals come to a river. A rowboat that seats two is available. If the cannibals ever outnumber the missionaries on either bank of the river, the missionaries will be eaten. (*AIMA problem 3.9*)

How shall they cross the river?



Formulation: Missionaries & Cannibals

- ***How to formalize:***

- *Initial state*: all M, all C, and boat on one bank
- *Actions*: ??
- *Transition Model*??
- *Goal test*: True if all M, all C, and boat on other bank
- *Cost*: ??

Remember:

- ***Representation:***

- *States*: Which properties matter & how to represent
- *Actions & Transition Model*: Which actions are possible & how to represent
- *Path Cost*: Deciding which action is next

Missionaries and Cannibals

States: (CL, ML, BL)

Initial 331

Goal 000

Actions:

Travel Across

Travel Back

-101

101

-201

201

-011

011

-021

021

-111

111

Outline for today's lecture

- Defining Task Environments
- Environment types
- Formulating Search Problems
- *Search Fundamentals (AIMA 3.3)*


Useful Concepts

- **State space:** the set of all states reachable from the initial state by *any* sequence of actions
 - *When several operators can apply to each state, this gets large very quickly*
 - *Might be a proper subset of the set of configurations*
- **Path:** a sequence of actions leading from one state s_j to another state s_k
- **Frontier:** those states that are available for *expanding* (for applying legal actions to)
- **Solution:** a path from the initial state s_i to a state s_f that satisfies the goal test

Basic search algorithms: *Tree Search*

- **Generalized algorithm to solve search problems (Review from CIS 121)**
 - *Enumerate in some order all possible paths from the initial state*
 - Here: search through *explicit tree generation*
 - ROOT= initial state.
 - Nodes in search tree generated through *transition model*
 - In general search generates a *graph* (same state through multiple paths), but we'll just look at *trees in lecture*
 - Tree search treats different paths to the same node as distinct

Review (CIS 121): Generalized tree search



function **TREE-SEARCH**(*problem*, *strategy*) return a solution or failure

 Initialize frontier to the *initial state* of the *problem*

 do

 if the frontier is empty then return *failure*

 choose leaf node for expansion according to *strategy* & remove from frontier


 if node contains goal state then return *solution*

 else expand the node and add resulting nodes to the frontier

Determines search
process!!

8-Puzzle: States and Nodes

- A **state** is a (representation of a) **physical configuration**
- A **node** is a data structure constituting **part of a search tree**
 - Also includes *parent, children, depth, path cost $g(x)$*
 - Here $node = \langle state, parent\text{-}node, children, action, path\text{-}cost, depth \rangle$
- **States do not have parents, children, depth or path cost!**

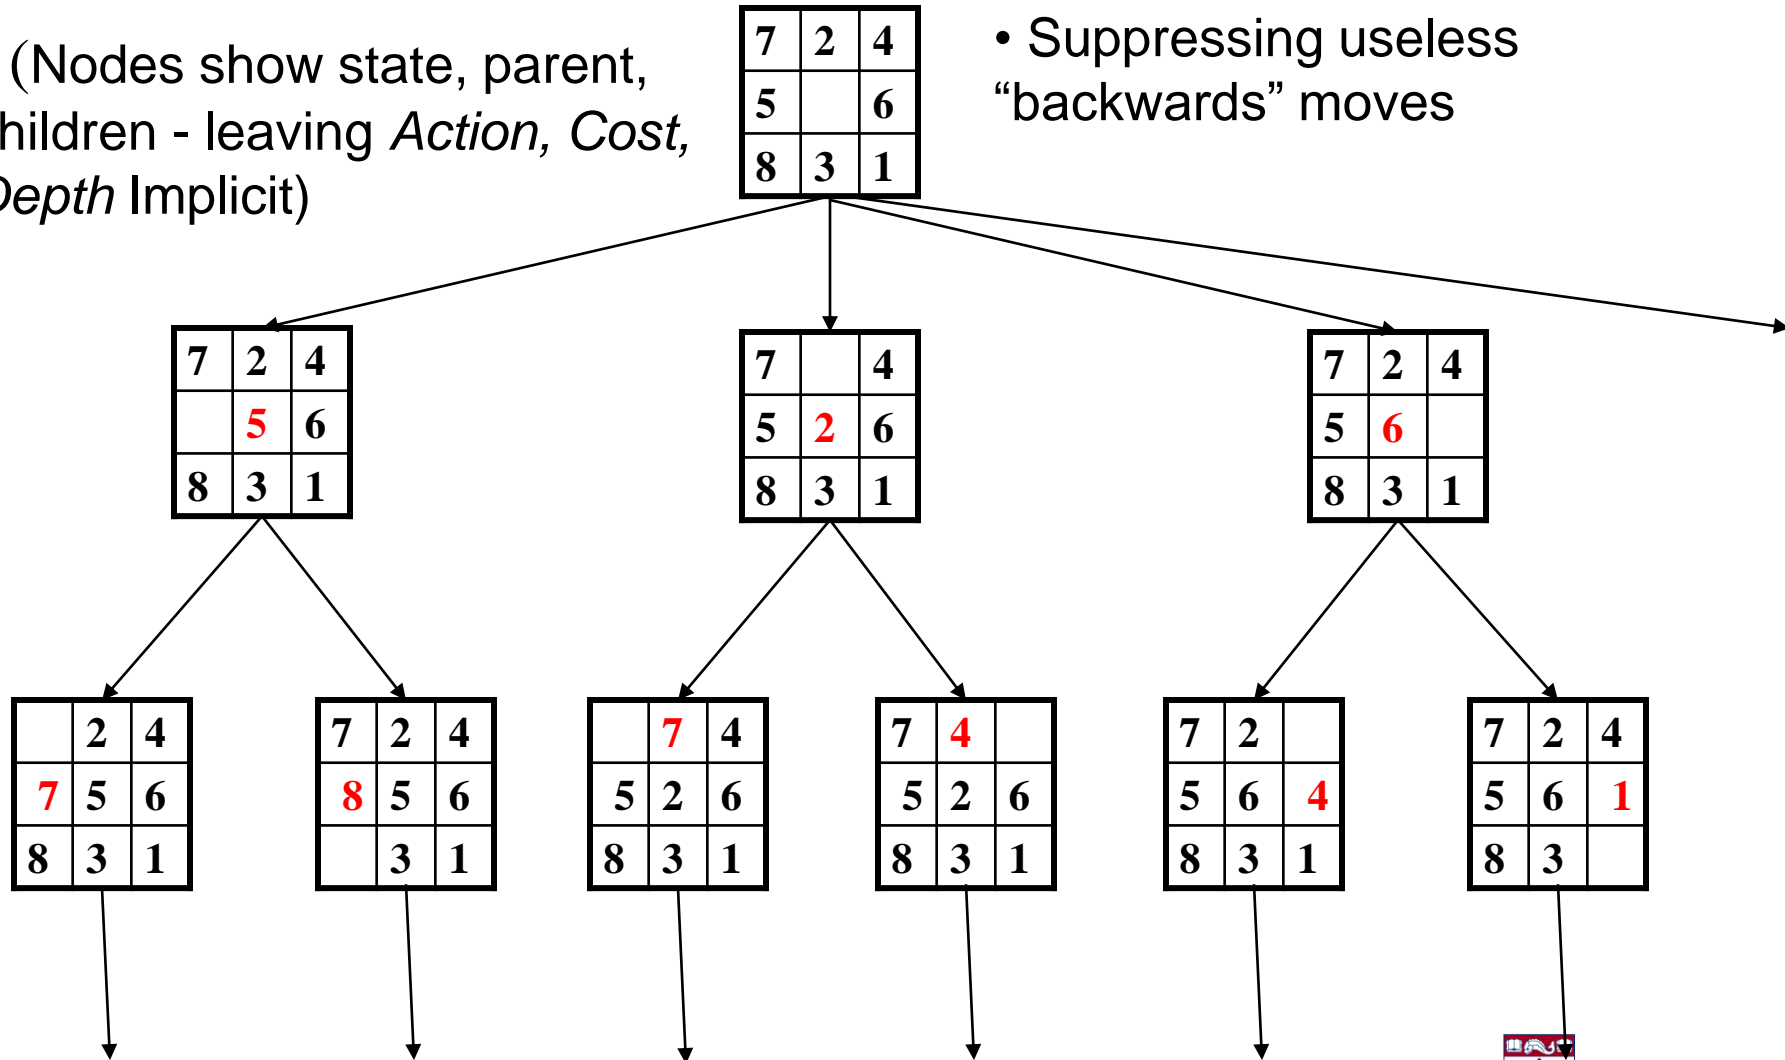


- **The EXPAND function**
 - uses the Actions and Transition Model to create the corresponding states
 - creates new nodes,
 - fills in the various fields

8-Puzzle *Search Tree*



- (Nodes show state, parent, children - leaving *Action*, *Cost*, *Depth* Implicit)

- Suppressing useless “backwards” moves




Problem: Repeated states


- Failure to detect *repeated states* can turn a linear problem into an *exponential* one!



Solution: Graph Search!



State Space



Search Tree

- **Graph search** ← *Optimal but memory inefficient*
 - Mod from tree search: Check to see if a node has been visited before adding to search queue
 - must keep track of all possible states (can use a lot of memory)
 - e.g., 8-puzzle problem, we have $9!/2 \approx 182K$ states

Graph Search vs Tree Search

function TREE-SEARCH(*problem*) **returns** a solution, or failure

initialize the frontier using the initial state of *problem*

loop do

if the frontier is empty **then return** failure

choose a leaf node and remove it from the frontier

if the node contains a goal state **then return** the corresponding solution

expand the chosen node, adding the resulting nodes to the frontier

function GRAPH-SEARCH(*problem*) **returns** a solution, or failure

initialize the frontier using the initial state of *problem*

initialize the explored set to be empty

loop do

if the frontier is empty **then return** failure

choose a leaf node and remove it from the frontier

if the node contains a goal state **then return** the corresponding solution

add the node to the explored set

expand the chosen node, adding the resulting nodes to the frontier

only if not in the frontier or explored set

Figure 3.7 An informal description of the general tree-search and graph-search algorithms. The parts of GRAPH-SEARCH marked in bold italic are the additions needed to handle repeated states.